

KHAJA RAHEEL AHMED MOHIUDDIN

Prompt Engineer | Adversarial LLM Evaluation | RLHF | LLM Systems | AI Operations

Warangal, India | +91-7013691213 | khajarahil0909@gmail.com | [Portfolio](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Prompt Engineer specializing in adversarial LLM evaluation, red-teaming, and advanced prompt design, building source-verified benchmark tasks that expose frontier-model web-search and retrieval failure modes. Previously contributed to Invisible Technologies' Meridial AI platform on Google-aligned projects, promoted to Reviewer after delivering 1500+ production-grade LLM training and evaluation tasks. Deep expertise across prompt design, RLHF pipelines, model alignment, and personalization-focused AI evaluation, backed by a software engineering background in Python, React, and production systems. Seeking to bring this combination of evaluation rigor and engineering fundamentals to a fast-moving product and engineering team.

CORE SKILLS

Adversarial Evaluation & Red-Teaming: Adversarial Testing / Red-Teaming, Model Stumping, Benchmark & Eval-Dataset Creation, Retrieval (RAG) Failure Analysis, Source Verification, Multi-Source Fact-Checking, Ground-Truth Annotation

AI Training & RLHF: LLM Training, RLHF Pipelines, Human Preference Data Curation, Response Ranking, Model Alignment

LLM Evaluation & Ops: Model Evaluation, Prompt Engineering, Adversarial Prompt Design, Multi-Hop Prompt Construction, Debugging LLM Outputs, Failure Analysis, Edge-Case Testing, AI Quality Review, Production AI Workflows

Domain Expertise: Python, Kotlin, JavaScript, TypeScript, React, REST APIs, Microservices

Tools & Platforms: Meridial AI Platform, Git, GitHub, GitLab, CI/CD, AWS SageMaker, SonarQube, Jest, Agile, Jira

PROFESSIONAL EXPERIENCE

Prompt Engineer - Adversarial LLM Evaluation & Red-Teaming

Jul 2026 – Present

Handshake AI

- Designed research-grade adversarial evaluation prompts ("model stumping") that probe web-search and retrieval-reasoning failures in a frontier large language model; every accepted task was required to defeat the model on ≥ 4 of 8 automated evaluation runs.
- Authored multi-hop, single-answer question, answer pairs with step-by-step "Golden Trajectory" solution paths, verifying each fact against 3+ independent primary sources for full reproducibility and zero-ambiguity grading.
- Diagnosed the core model failure mode, genuine information un-retrievability vs. reasoning difficulty and codified a reusable prompt-construction methodology that lifted first-pass acceptance and throughput across multiple seeded domains.
- Ran controlled model evaluations (memory-off, extended-reasoning, 3-run + server-side 8-run auto-eval) and wrote precise failure justifications separating retrieval failures from ambiguity, calculation, or prompt error.
- Drove tasks through a multi-stage QA pipeline (auto-eval $\rightarrow$ two human review rounds $\rightarrow$ audit), resolving granular reviewer feedback on prompt ambiguity, source validity, and source-to-trajectory mapping to clean approval.
- Delivered multiple approved evaluation tasks across a wide range of domains, including Art, Music, Geography, Education, Science & Technology, and more, to scale the method.

Master - Python Coding Specialist - AI Trainer

Oct 2025 – Jun 2026

Meridial AI – Invisible Technologies (Google-aligned Projects)

- Delivering 1500+ high-quality LLM training and evaluation tasks on Invisible Technologies' Meridial AI platform, directly improving model accuracy, reasoning, and instruction-following behavior for Google-aligned AI systems.
- Promoted to Reviewer, responsible for auditing 1000+ model outputs for logical correctness, edge-case handling, and production-grade quality standards, ensuring Invisible's AI training output meets enterprise client expectations.
- Conducted deep LLM performance analysis, identifying failure patterns and implementing structured human feedback loops to improve response quality and reduce model errors at scale.
- Led evaluation of code-generation models (Python, Kotlin), validating output correctness and adherence to software engineering best practices.
- Applied RLHF methodologies to strengthen model alignment and robustness, supporting production-grade results for enterprise clients.

AI Training Data Specialist

Mar 2025 – Sep 2025

Soul AI (Deccan AI)

- Contributed to RLHF-based LLM training pipelines, curating preference data and improving model instruction-following, contextual reasoning, and output consistency at scale.
- Evaluated and refined LLM responses across diverse task categories, applying production-readiness benchmarks to ensure model reliability and quality.

- Identified recurring failure patterns in model outputs and provided corrective annotation signals, directly improving training data quality.

Associate Software Engineer

Dec 2023 – Feb 2025

Tech Mahindra

- Built and optimized production-grade frontend systems for Singapore Land Authority, improving application performance, scalability, and user experience.
- Collaborated with backend teams on API integration and distributed systems workflows; implemented unit tests using Jest to increase code quality and reduce production defects.
- Played a key role in a successful December 2024 production deployment, contributing to design and code reviews that ensured best practices across the team.

KEY PROJECTS

Project: Seal (Frontier-Model Web-Search Benchmark)

2026

Handshake AI

- Built adversarial, source-verified evaluation prompts engineered to make a frontier LLM fail its web-search benchmark ($\geq 4/8$ auto-eval), each backed by a reproducible multi-source solution trajectory and iterated through a multi-stage human-review pipeline.

Project: Grizzly (Google - LLM Code Intelligence Simulation)

2025

Meridial AI – Invisible Technologies

- Simulated real-world developer workflows using LLM systems; resolved repo-level bug reports and feature requests, delivering production-grade Kotlin patches with testing and validation.

Project: Toucan (Google - RLHF System Training)

2025

Meridial AI – Invisible Technologies

- Delivered structured human feedback to improve LLM response quality across code correctness, logical reasoning, and instruction-following, strengthening model alignment and consistency.

Project: P13N Personalization (Google - AI Personalization System)

2025

Meridial AI – Invisible Technologies

- Evaluated context-aware AI response generation by comparing side-by-side model outputs with and without personal context integrated, using historical interaction data to support model training decisions.

RAG Evaluation Service (Production-Style LLM Microservice)

2026

LLM Systems Project

- Built a production-style Retrieval-Augmented Generation microservice with hybrid retrieval (BM25 + TF-IDF, Reciprocal Rank Fusion), cross-encoder re-ranking, versioned prompts, PII-redaction guardrails, and a human-in-the-loop gate (FastAPI + LangChain).
- Designed a reproducible evaluation harness measuring retrieval quality (hit-rate, MRR) and answer grounding, with the model refusing rather than hallucinating when no relevant context is found.

IoT Intrusion Detection System (PB-DID)

2023

Deep Learning Research

- Built a stacked-LSTM deep learning classifier achieving 98.05% accuracy on 42,482 held-out network flows (Normal / DoS / DDoS) using UNSW-NB15 and Bot-IoT protocol features, served via a Flask dashboard, CLI, and JSON API.

Airfare Prediction System

2023

Machine Learning Project

- Built an end-to-end stacked-ensemble regression model (LightGBM, XGBoost, CatBoost, Extra Trees) reaching 94.14% accuracy (R^2 0.9165) on 10,683 itineraries under 10-fold cross-validation, powering cost-effective booking recommendations.

EDUCATION

B.Tech in Computer Science & Engineering

2019 – 2023

Balaji Institute of Technology and Science, Warangal | CGPA: 7.71

CERTIFICATIONS & ACHIEVEMENTS

- IEEE Publication – Evolution of Data Science in IT · Cisco – Python & Cybersecurity Essentials
- AWS SageMaker (No-Code ML) & Streamlit · Virtual Experience Programs – Deloitte, Goldman Sachs, JP Morgan